

In []:

```
# references:
# https://github.com/Kaggle/kaggle-api
# https://towardsdatascience.com/downloading-datasets-into-google-drive-via-google-colab-bcblb30b0166
# https://www.pyimagesearch.com/2018/12/24/how-to-use-keras-fit-and-fit_generator-a-hands-on-tutorial/
```

In [1]:

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

In [2]:

```
# importing necessary libraries
import tensorflow as tf
import datetime
import os
from tensorflow.keras.layers import Dense, Activation, Conv2D, Flatten, MaxPooling2D, Dropout
from tensorflow.keras import regularizers, optimizers, initializers
from tensorflow.keras.models import Model
from keras.preprocessing.image import ImageDataGenerator
import pandas as pd
import numpy as np
```

In [4]:

```
!pip install -q kaggle
from google.colab import files
files.upload()
!mkdir ~/.kaggle
!cp kaggle.json ~/.kaggle/
! chmod 600 ~/.kaggle/kaggle.json
```

Choose File No file selected

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

In [5]:

```
!kaggle datasets download -d brahma0545/aaic-assignment-tl
```

```
Downloading aaic-assignment-tl.zip to /content
100% 4.33G/4.34G [01:24<00:00, 67.3MB/s]
100% 4.34G/4.34G [01:24<00:00, 55.2MB/s]
```

In [6]:

```
# list the files in current directory:
!ls -ltr
```

```
total 4549336
drwxr-xr-x 1 root root      4096 Feb  4 15:26 sample_data
drwx----- 5 root root      4096 Feb  8 04:56 gdrive
-rw-r--r-- 1 root root         62 Feb  8 05:11 kaggle.json
-rw-r--r-- 1 root root 4658501240 Feb  8 05:12 aaic-assignment-tl.zip
```

In []:

```
!unzip "/content/aaic-assignment-tl.zip" -d "/content/TL"
```

In [8]:

```
# label data file
dir_path = "/content/TL/labels_final.csv"
```

In [9]:

```
# fetch labels_final.csv
train_df = pd.read_csv(dir_path)
```

In [10]:

```
train_df.head()
```

Out[10]:

	path	label
0	imagesv\o\h\voh71d00\509132755+-2755.tif	3
1	imagesl\l\xt\xt19d00\502213303.tif	3
2	imagesx\xe\d\xed05a00\2075325674.tif	2
3	imageso\o\j\b\obj60d00\517511301+-1301.tif	3
4	imagesq\q\z\k\qzk17e00\2031320195.tif	7

In [11]:

```
# replacing labels
train_df = train_df.replace({'label':
                             {0:"letter",
                              1:"form",
                              2:"email",
                              3:"handwritten",
                              4:"advertisement",
                              5:"scientifit report",
                              6:"scientific publication",
                              7:"specification",
                              8:"file folder",
                              9:"news article",
                              10:"budget",
                              11:"invoice",
                              12:"presentation",
                              13:"questionnaire",
                              14:"resume",
                              15:"memo"}}})
```

In [12]:

```
# how much data for each category:
train_df['label'].value_counts()
```

Out[12]:

letter	3016
questionnaire	3007
presentation	3006
resume	3006
handwritten	3005
file folder	3003
budget	3002
news article	3002
specification	3000
scientifit report	2999

```
memo                2996
form                2994
advertisement       2994
email               2993
invoice             2992
scientific publication 2985
Name: label, dtype: int64
```

Observations:

1. data is balanced

In [13]:

```
datagen = ImageDataGenerator(rotation_range=90,width_shift_range=0.25,height_shift_range=0.25,horizontal_flip=0.25,vertical_flip=0.25,
                             rescale=1./255,validation_split=0.30)
```

In [14]:

```
# train data generator
train_generator = datagen.flow_from_dataframe(
    dataframe=train_df,
    directory="/content/TL/data_final/",
    x_col = "path",
    y_col = "label",
    subset = "training",
    batch_size = 96,
    seed = 39,
    shuffle = True,
    class_mode="categorical",
    target_size = (224,224)
)
```

Found 33600 validated image filenames belonging to 16 classes.

In [15]:

```
# validation data generator
valid_generator = datagen.flow_from_dataframe(
    dataframe=train_df,
    directory="/content/TL/data_final/",
    x_col = "path",
    y_col = "label",
    subset = "validation",
    batch_size = 32,
    seed = 42,
    shuffle = True,
    class_mode="categorical",
    target_size = (224,224)
)
```

Found 14400 validated image filenames belonging to 16 classes.

In [16]:

```
%load_ext tensorboard
import tensorflow as tf
import datetime, os
```

VGG16 Pretrained model as base model

Model-3

In [17]:

```
tf.keras.backend.clear_session()

# loading vgg16 from keras
from keras.applications import VGG16
# load model
base_model = VGG16(include_top=False, input_shape = (224, 224, 3), weights='imagenet')
# need to train last 6 layers of VGG-16 network
for i, layer in enumerate(base_model.layers):
    if (i < 13):
        layer.trainable = False
    else :
        layer.trainable = True

# Conv layer
conv1 = Conv2D(512, kernel_size=(7, 7), padding='valid', strides=(1, 1), activation='relu')(base_model.output)

# Conv layer
conv2 = Conv2D(128, kernel_size=(1, 1), padding='valid', strides=(1, 1), activation='relu')(conv1)

# Flatten
flat = Flatten()(conv2)

# output layer
model3 = Dense(16, activation='softmax')(flat)
final_model3 = Model(base_model.input, model3)
final_model3.compile(optimizer= optimizers.SGD(learning_rate=0.001, momentum=0.9),
                    loss="categorical_crossentropy",
                    metrics = ['accuracy'])
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58892288/58889256 [=====] - 0s 0us/step

In [22]:

```
final_model3.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0

block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 1, 1, 512)	12845568
conv2d_1 (Conv2D)	(None, 1, 1, 128)	65664
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 16)	2064
=====		
Total params: 27,627,984		
Trainable params: 22,352,528		
Non-trainable params: 5,275,456		

In [18]:

```
!rm -rf ./logs/
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=1)
vgghist = final_model3.fit(train_generator,
                           validation_data=valid_generator,
                           callbacks=tensorboard_callback,
                           steps_per_epoch = 350,
                           epochs = 15
                           )
```

```
Epoch 1/15
350/350 [=====] - 753s 2s/step - loss: 2.2462 - accuracy: 0.2739 - val_loss: 1
.5777 - val_accuracy: 0.5018
Epoch 2/15
350/350 [=====] - 734s 2s/step - loss: 1.5491 - accuracy: 0.5140 - val_loss: 1
.4659 - val_accuracy: 0.5381
Epoch 3/15
350/350 [=====] - 713s 2s/step - loss: 1.4142 - accuracy: 0.5552 - val_loss: 1
.3634 - val_accuracy: 0.5717
Epoch 4/15
350/350 [=====] - 687s 2s/step - loss: 1.3296 - accuracy: 0.5836 - val_loss: 1
.3135 - val_accuracy: 0.5902
Epoch 5/15
350/350 [=====] - 679s 2s/step - loss: 1.2846 - accuracy: 0.5990 - val_loss: 1
.2661 - val_accuracy: 0.6032
Epoch 6/15
350/350 [=====] - 668s 2s/step - loss: 1.2540 - accuracy: 0.6070 - val_loss: 1
.2783 - val_accuracy: 0.5998
Epoch 7/15
350/350 [=====] - 669s 2s/step - loss: 1.2197 - accuracy: 0.6182 - val_loss: 1
.2207 - val_accuracy: 0.6212
Epoch 8/15
350/350 [=====] - 668s 2s/step - loss: 1.1798 - accuracy: 0.6309 - val_loss: 1
.2176 - val_accuracy: 0.6245
Epoch 9/15
350/350 [=====] - 673s 2s/step - loss: 1.1514 - accuracy: 0.6432 - val_loss: 1
.1881 - val_accuracy: 0.6338
Epoch 10/15
350/350 [=====] - 672s 2s/step - loss: 1.1076 - accuracy: 0.6531 - val_loss: 1
.1573 - val_accuracy: 0.6435
Epoch 11/15
350/350 [=====] - 672s 2s/step - loss: 1.0949 - accuracy: 0.6598 - val_loss: 1
.1255 - val_accuracy: 0.6522
Epoch 12/15
350/350 [=====] - 672s 2s/step - loss: 1.0574 - accuracy: 0.6690 - val_loss: 1
.1192 - val_accuracy: 0.6544
Epoch 13/15
350/350 [=====] - 676s 2s/step - loss: 1.0420 - accuracy: 0.6725 - val_loss: 1
.1282 - val_accuracy: 0.6524
Epoch 14/15
350/350 [=====] - 676s 2s/step - loss: 1.0284 - accuracy: 0.6787 - val_loss: 1
```

.0977 - val_accuracy: 0.6619

Epoch 15/15

350/350 [=====] - 679s 2s/step - loss: 1.0170 - accuracy: 0.6837 - val_loss: 1

.0660 - val_accuracy: 0.6708

In []:

```
# !tensorboard dev upload --logdir /content/logs \  
#   --name "Transfer Learning Model:3 (15 epochs)" \  
#   --description " from TF_3.ipynb " \  
#   --one_shot
```

In []: